

## Lab 2 Exercise – PyTorch AutoGrad

Justin Ugwudike (jknu1g19@soton.ac.uk)

Feb 27, 2023

### 1.1 Question 1\_1

$$\min_{\hat{U}, \hat{V}} (\|A - \hat{U}\hat{V}^T\|_F^2)$$

Figure 1 – The function to optimise.

Using the code shown in Appendix A, the rank-2 factorisation of the matrix is achieved by minimising the function shown in Figure 1. The results are shown in Figure 2.

```
RECONSTRUCTION-----  
tensor([[ 0.2193,  0.5037,  0.3697],  
        [ 3.2610, -0.0122,  1.9618],  
        [ 3.0304,  0.6037,  2.1129]])  
MSE LOSS-----  
tensor(0.1224)
```

Figure 2 – The reconstruction of the target matrix using U and V (obtained via gradient descent utilising automatic differentiation). The mean squared error loss is shown.

The code produces an approximation of the target matrix with the MSE loss being 0.1224, this value differs by +0.0004 to the loss achieved using stochastic gradient descent with non-automatic gradients in Lab 1.

### 1.2 Question 1\_2

With the code from Appendix A the rank-2 factorisation of the data from a dataset of 150 instances and 4 features was estimated using gradient descent. The function in Appendix B allowed me to reconstruct the data matrix using Singular Value decomposition (SVD), the losses achieved using both methods is shown in Figure 3.

```
MSE LOSS-----  
tensor(15.2292, dtype=torch.float64)  
SVD MSE LOSS-----  
tensor(15.2288, dtype=torch.float64)
```

Figure 3 – The reconstruction loss of the data matrix through gradient descent and the rank-2 reconstruction loss through SVD.

The reconstruction loss of the data matrix using the function in Appendix A is 15.2292. The losses achieved when using each method are very close in value (0.0004 difference).

### 1.3 Question 1\_3




Figure 4 – The first graph represents the data in the U matrix computed by SVD and the second shows the data computed by gradient descent.

By looking at Figure 4 both methods have grouped the data into two classes. The two pairs of classes retain similar shapes and sizes when using gradient descent and SVD however, the orientations of the two pairs of data groups are different.

### 1.4 Question 2\_1

The function used to train an MLP through gradient descent is shown in Appendix C.

### 1.5 Question 2\_2

From Appendix D I can see that either increasing the learning rate and maximum iterations leads to a higher prediction accuracy in both the training and validation data. In every experiment, the accuracy when using the training data is much greater than the accuracy when using the validation data. From these results I can infer that using an iteration count of 1000 and learning rate of 0.01 (to avoid overfitting when using a learning rate of 0.1) would be the best parameters to use in this example as the accuracies for both the training and validation data reach 99% and above.

## 2 Appendix

Appendix A - Completed `gd_factorise_ad()` function.

```
20 def gd_factorise_ad(A: torch.Tensor, rank : int, num_epochs=1000, lr=0.01):
21     U_init = torch.rand(A.shape[0], rank)
22     V_init = torch.rand(A.shape[1], rank)
23     U = torch.tensor(U_init, requires_grad=True, dtype=torch.float)
24     V = torch.tensor(V_init, requires_grad=True, dtype=torch.float)
25
26
27     for i in range(0, num_epochs):
28         # manually dispose of the gradient (in reality it would be better to det
29         U.grad=None
30         V.grad=None
31         # evaluate the function
32         J = function(A,U,V)
33         # auto-compute the gradients at the previously evaluated point x
34         J.backward()
35         # compute the update
36         U.data = U - U.grad*lr
37         V.data = V - V.grad*lr
38
39     return U, V
```

Appendix B – Singular Value Decomposition and reconstruction function.

```
67 def SVD_reonstrucion(A):
68     SVD = torch.svd(A) #svd
69     singular_values = torch.diag(SVD.S) #singular values
70     #print("SINGULAR VALUES-----")
71     #print(singular_values)
72     singular_values[singular_values.shape[0] - 1] = 0 #set last singular value
73     singular_values[singular_values.shape[0] - 2] = 0 #set second-to-last singu
74     #print("RANK2 SINGULAR VALUES-")
75     #print(singular_values)
76     reconstruction = SVD.U @ singular_values @ SVD.V.t()
77     mse_loss = torch.nn.functional.mse_loss(reconstruction, A, reduction='sum')
78     return reconstruction, mse_loss
```

Appendix C – Function to perform MLP gradient descent using automatic differentiation.

```
145 def MLP(data_in, targets_in, num_epochs=100, lr=0.01):
146     W1_init = torch.rand(4, 12)
147     W2_init = torch.rand(12, 3)
148     b_init = torch.tensor([0])
149     W1 = torch.tensor(W1_init, requires_grad=True, dtype=torch.float)
150     W2 = torch.tensor(W2_init, requires_grad=True, dtype=torch.float)
151     b1 = torch.tensor(b_init, requires_grad=True, dtype=torch.float)
152     b2 = torch.tensor(b_init, requires_grad=True, dtype=torch.float)
153
154     for i in range(0, num_epochs):
155         # manually dispose of the gradient (in reality it would be better to de
156         W1.grad=None
157         W2.grad=None
158         b1.grad=None
159         b2.grad=None
160         # evaluate the function
161         logits = torch.relu(data_in @ W1 + b1) @ W2 + b2
162         J = torch.nn.functional.cross_entropy(logits, targets_in)
163         # auto-compute the gradients at the previously evaluated point x
164         J.backward()
165         # compute the update
166         W1.data = W1 - W1.grad*lr
167         W2.data = W2 - W2.grad*lr
168         b1.data = b1 - b1.grad*lr
169         b2.data = b2 - b2.grad*lr
170
171     return W1,W2,b1,b2,J
```

Appendix D – Three sets of accuracies recorded from an MLP when using different learning rates and maximum iterations.

ACCURACY-----		ACCURACY-----		ACCURACY-----	
LR=0.1	ITER=100	LR=0.1	ITER=100	LR=0.1	ITER=100
training	99.0%	training	99.0%	training	99.0%
validation	100.0%	validation	100.0%	validation	100.0%
LR=0.01	ITER=100	LR=0.01	ITER=100	LR=0.01	ITER=100
training	69.0%	training	79.0%	training	79.0%
validation	100.0%	validation	100.0%	validation	100.0%
LR=0.001	ITER=100	LR=0.001	ITER=100	LR=0.001	ITER=100
training	51.0%	training	41.0%	training	48.0%
validation	72.0%	validation	98.0%	validation	72.0%
LR=0.1	ITER=1000	LR=0.1	ITER=1000	LR=0.1	ITER=1000
training	99.0%	training	100.0%	training	100.0%
validation	100.0%	validation	100.0%	validation	100.0%
LR=0.01	ITER=1000	LR=0.01	ITER=1000	LR=0.01	ITER=1000
training	99.0%	training	99.0%	training	99.0%
validation	100.0%	validation	100.0%	validation	100.0%
LR=0.001	ITER=1000	LR=0.001	ITER=1000	LR=0.001	ITER=1000
training	94.0%	training	80.0%	training	84.0%
validation	98.0%	validation	100.0%	validation	98.0%